

Snellius File Systems

Hands-on data management in high-performance computing

Xavier Álvarez-Farré

High-Performance Computing and Visualization, SURF, Science Park 140, 1098XG Amsterdam, The Netherlands

October 9, 2024



Introduction

HPC applications are software designed to exploit the power of supercomputers or high-performance clusters to solve complex problems, perform intensive **computations**, facilitate **communications** between parallel processes, and manage **input/output** operations on a high-performance file system.

HPC applications are software designed to exploit the power of supercomputers or high-performance clusters to solve complex problems, perform intensive **computations**, facilitate **communications** between parallel processes, and manage **input/output** operations on a high-performance file system.

Computations

Extensive numerical operations on a specific region of interest (*e.g.*, mesh elements, particles, or even words in a text document).

- Scale $\mathcal{O}(P)$ for P processors.

HPC applications are software designed to exploit the power of supercomputers or high-performance clusters to solve complex problems, perform intensive **computations**, facilitate **communications** between parallel processes, and manage **input/output** operations on a high-performance file system.

Computations

Extensive numerical operations on a specific region of interest (e.g., mesh elements, particles, or even words in a text document).

- Scale $\mathcal{O}(P)$ for P processors.

Communications

Coordination and data exchanges between parallel processes to produce meaningful results.

- Scale $\mathcal{O}(\sqrt{P})$ for P processors.

HPC applications are software designed to exploit the power of supercomputers or high-performance clusters to solve complex problems, perform intensive **computations**, facilitate **communications** between parallel processes, and manage **input/output** operations on a high-performance file system.

Computations

Extensive numerical operations on a specific region of interest (e.g., mesh elements, particles, or even words in a text document).

- Scale $\mathcal{O}(P)$ for P processors.

Communications

Coordination and data exchanges between parallel processes to produce meaningful results.

- Scale $\mathcal{O}(\sqrt{P})$ for P processors.

I/O operations

Interactions with the file system to read input data, write output results, and perform intermediate data storage.

- Usually have no scaling factor.

Input operations

- **Initialization data** such as geometric grids, genomic sequences, climate simulation data, or large text files.
 - Massive files that require a high level of precision.
 - Involve parallel input operations.

Input operations

- **Initialization data** such as geometric grids, genomic sequences, climate simulation data, or large text files.
 - Massive files that require a high level of precision.
 - Involve parallel input operations.
- **Job parameters** that provide instructions for guiding the behavior of the job and coordinating processes are typically stored in small files but require coordination between processes.

Input operations

- **Initialization data** such as geometric grids, genomic sequences, climate simulation data, or large text files.
 - Massive files that require a high level of precision.
 - Involve parallel input operations.
- **Job parameters** that provide instructions for guiding the behavior of the job and coordinating processes are typically stored in small files but require coordination between processes.

Output operations

- **Restart files** or checkpoints enable job continuation in the event of an interruption.
 - Large files that require full precision of simulation data.
 - Involve periodic parallel output operations.

Input operations

- **Initialization data** such as geometric grids, genomic sequences, climate simulation data, or large text files.
 - Massive files that require a high level of precision.
 - Involve parallel input operations.
- **Job parameters** that provide instructions for guiding the behavior of the job and coordinating processes are typically stored in small files but require coordination between processes.

Output operations

- **Restart files** or checkpoints enable job continuation in the event of an interruption.
 - Large files that require full precision of simulation data.
 - Involve periodic parallel output operations.
- **Analysis data** for visualization, post-processing and in-depth examination.
 - May be possible to use reduced precision or store a subset of the full-resolution data.
 - Involve periodic parallel output operations.


Input operations

- **Initialization data** such as geometric grids, genomic sequences, climate simulation data, or large text files.
 - Massive files that require a high level of precision.
 - Involve parallel input operations.
- **Job parameters** that provide instructions for guiding the behavior of the job and coordinating processes are typically stored in small files but require coordination between processes.

Output operations

- **Restart files** or checkpoints enable job continuation in the event of an interruption.
 - Large files that require full precision of simulation data.
 - Involve periodic parallel output operations.
- **Analysis data** for visualization, post-processing and in-depth examination.
 - May be possible to use reduced precision or store a subset of the full-resolution data.
 - Involve periodic parallel output operations.
- **Job monitoring** data that provides insights into the application's performance and progress is typically generated by process 0 in a human-readable format, such as ASCII.

Data management plays a crucial role in HPC applications, encompassing **diverse tasks** with varying natures and specific requirements. At the same time, HPC systems are equipped with **distinct file systems** to meet these specific needs.



Data management plays a crucial role in HPC applications, encompassing **diverse tasks** with varying natures and specific requirements. At the same time, HPC systems are equipped with **distinct file systems** to meet these specific needs.

File-system	Speed	Size	Backed-up	Expiration
/home	Normal	200 GB	Yes	Permanent
/scratch-shared	Fast		No	6/14 days
/scratch-local	Fast	8 TB	No	6/14 days
/scratch-node	Fastest		No	End of job
/project	Fast	On demand	No	Permanent
/archive	Slow	On demand	Yes	Permanent

Hands-on

1. Navigating file systems:

- Understand the structure of HPC file systems (home, project, scratch, archive)
- Move between directories and check the current location

2. File management:

- Create, view, and edit files using basic terminal commands
- Move, rename, and copy files between different file systems

3. Directory management:

- Create and delete directories
- Navigate directories and manage file hierarchy

4. File permissions:

- Check and modify file access permissions

5. Disk usage monitoring:

- Monitor storage usage and available space across file systems

Objective: Learn how to move between directories in different file systems (home, project, scratch, archive) and understand the structure of the HPC environment.

Key commands:

- `pwd` - Display the current directory
- `ls` - List the contents of a directory
- `cd <path>` - Change directories

Tasks:

- Start in the **home** directory and use `pwd` to check your location.
- Use `cd` to navigate to **scratch** directory.
- Explore the contents of **scratch** directory with `ls`.
- Return to the **home** directory using `cd`.

Expected output:

- Comfortable navigating and exploring directories in the HPC file system.

Objective: Learn how to create, view, move, and rename files across different file systems in an HPC environment.

Key commands:

- `touch <fname>` - Create a new empty file
- `nano <fname>` or `vim <filename>` - Edit a file
- `cat <fname>` - Display the contents of a file
- `cp <src> <tgt>` - Copy files between directories
- `mv <src> <tgt>` - Move or rename files

Tasks:

- Create a file in the **home** directory using **touch**.
- Edit the file with **nano** or **vi**, and view its contents with **cat**.
- Copy the file to the **scratch** directory using **cp**.
- Rename or move the file using **mv**.

Expected output:

- Understand how to manage files: creation, editing, copying, and renaming.

Objective: Learn how to create, navigate, and remove directories within the HPC file system.

Key commands:

- `mkdir <dname>` - create a new directory
- `cd <dname>` - navigate into a directory
- `rmdir <dname>` - remove an empty directory
- `ls [dname]` - list directory contents

Tasks:

- Create a new directory in the home directory using `mkdir`.
- Navigate into the new directory with `cd`.
- List its contents with `ls`, then return to the home directory.
- Remove the directory using `rmdir`.

Expected output:

- Understand how to create, enter, and remove directories and manage directory structure effectively.

Objective: Learn how to view and modify file access permissions to control who can read, write, or execute files.

Key commands:

- `ls -l` - view file permissions
- `chmod` - change file permissions

Explanation:

- Permissions are shown as `rw-rw-rw-` (for owner, group, others).
- `r` - read (4), `w` - write (2), `x` - execute (1).
- Numeric format: `chmod 755` means owner can `rw-`, group and others can `r-x`.

Tasks:

- Use `ls -l` to check the permissions of a file.
- Modify the file's permissions to allow read access for everyone using `chmod 644`.
- Verify the permission changes with `ls -l`.

Expected output:

- Understand how to check and modify file permissions in the HPC environment.

Objective: Learn how to monitor storage usage and check available space across different file systems.

Key commands:

- `df -h` - view disk space usage for all mounted file systems
- `du -sh` - check the size of a specific directory

Tasks:

- Use `df -h` to view overall disk usage and available space in the home, scratch, and project file systems.
- Check the size of a specific directory in your home using `du -sh`.
- Identify where you are using the most space.

Expected output:

- Understand how to monitor storage usage and manage disk space efficiently.