Snellius File Systems

Hands-on data management in high-performance computing

Xavier Álvarez-Farré
High-Performance Computing and Visualization, SURF, Science Park 140, 1098XG Amsterdam, The Netherlands

October 29, 2023

SURF

EURO NETHERLANDS

Introduction

Overview of HPC applications

HPC applications are software designed to exploit the power of supercomputers or high-performance clusters to solve complex problems, perform intensive computations, facilitate communications between parallel processes, and manage input/output operations on a high-performance file system.

Overview of HPC applications

HPC applications are software designed to exploit the power of supercomputers or high-performance clusters to solve complex problems, perform intensive computations, facilitate communications between parallel processes, and manage input/output operations on a high-performance file system.

Computations

Extensive numerical operations on a specific region of interest (*e.g.*, mesh elements, particles, or even words in a text document).

- Scale $\mathcal{O}(P)$ for P processors.

Overview of HPC applications

HPC applications are software designed to exploit the power of supercomputers or high-performance clusters to solve complex problems, perform intensive computations, facilitate communications between parallel processes, and manage input/output operations on a high-performance file system.

Computations

Extensive numerical operations on a specific region of interest (*e.g.*, mesh elements, particles, or even words in a text document).

- Scale $\mathcal{O}(P)$ for P processors.

Communications

Coordination and data exchanges between parallel processes to produce meaningful results.

- Scale $\mathcal{O}(\sqrt{P})$ for P processors.

Overview of HPC applications

HPC applications are software designed to exploit the power of supercomputers or high-performance clusters to solve complex problems, perform intensive computations, facilitate communications between parallel processes, and manage input/output operations on a high-performance file system.

Computations

Extensive numerical operations on a specific region of interest (*e.g.*, mesh elements, particles, or even words in a text document).

- Scale $\mathcal{O}(P)$ for P processors.

Communications

Coordination and data exchanges between parallel processes to produce meaningful results.

- Scale $\mathcal{O}(\sqrt{P})$ for P processors.

I/O operations

Interactions with the file system to read input data, write output results, and perform intermediate data storage.

- Usually have no scaling factor.

Input operations

- **Initialization data** such as geometric grids, genomic sequences, climate simulation data, or large text files.
  - Massive files that require a high level of precision.
  - Involve parallel input operations.

Input operations

- **Initialization data** such as geometric grids, genomic sequences, climate simulation data, or large text files.
  - Massive files that require a high level of precision.
  - Involve parallel input operations.
- **Job parameters** that provide instructions for guiding the behavior of the job and coordinating processes are typically stored in small files but require coordination between processes.

Overview of I/O operations

Input operations

- **Initialization data** such as geometric grids, genomic sequences, climate simulation data, or large text files.
  - Massive files that require a high level of precision.
  - Involve parallel input operations.
- **Job parameters** that provide instructions for guiding the behavior of the job and coordinating processes are typically stored in small files but require coordination between processes.

Output operations

- **Restart files** or checkpoints enable job continuation in the event of an interruption.
  - Large files that require full precision of simulation data.
  - Involve periodic parallel output operations.

Overview of I/O operations

Input operations

- **Initialization data** such as geometric grids, genomic sequences, climate simulation data, or large text files.
  - Massive files that require a high level of precision.
  - Involve parallel input operations.
- **Job parameters** that provide instructions for guiding the behavior of the job and coordinating processes are typically stored in small files but require coordination between processes.

Output operations

- **Restart files** or checkpoints enable job continuation in the event of an interruption.
  - Large files that require full precision of simulation data.
  - Involve periodic parallel output operations.
- **Analysis data** for visualization, post-processing and in-depth examination.
  - May be possible to use reduced precision or store a subset of the full-resolution data.
  - Involve periodic parallel output operations.

Overview of I/O operations

Input operations

- **Initialization data** such as geometric grids, genomic sequences, climate simulation data, or large text files.
  - Massive files that require a high level of precision.
  - Involve parallel input operations.
- **Job parameters** that provide instructions for guiding the behavior of the job and coordinating processes are typically stored in small files but require coordination between processes.

Output operations

- **Restart files** or checkpoints enable job continuation in the event of an interruption.
  - Large files that require full precision of simulation data.
  - Involve periodic parallel output operations.
- **Analysis data** for visualization, post-processing and in-depth examination.
  - May be possible to use reduced precision or store a subset of the full-resolution data.
  - Involve periodic parallel output operations.
- **Job monitoring** data that provides insights into the application's performance and progress is typically generated by process 0 in a human-readable format, such as ASCII.

Overview of Snellius file-systems

Data management plays a crucial role in HPC applications, encompassing diverse tasks with varying natures and specific requirements. At the same time, HPC systems are equipped with distinct file systems to meet these specific needs.

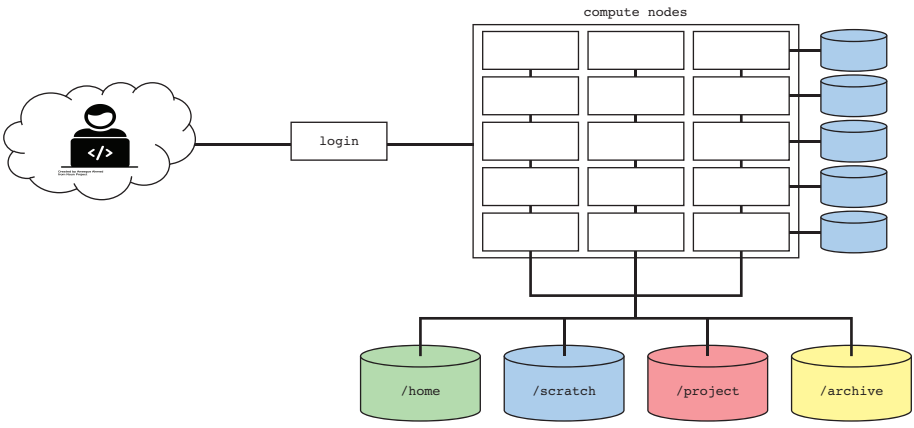

Overview of Snellius file-systems

Data management plays a crucial role in HPC applications, encompassing diverse tasks with varying natures and specific requirements. At the same time, HPC systems are equipped with distinct file systems to meet these specific needs.

File-system	Speed	Size	Backed-up	Expiration
`/home`	Normal	200 GB	Yes	Permanent
`/scratch-shared`	Fast		No	6/14 days
`/scratch-local`	Fast	8 TB	No	6/14 days
`/scratch-node`	Fastest		No	End of job
`/project`	Fast	On demand	No	Permanent
`/archive`	Slow	On demand	Yes	Permanent

Hands-on